



South Carolina


STAKEHOLDER SUMMARY

Bot detection and gift-card decisions

Which caregiver responses are real, which need a second look, and who should be paid.

1,956 responses reviewed

Caregiver Acceptability Study · Infant Autism Screening · September 2026



Two study databases, one shared check

Every response in both databases went through the same screen, so the two are directly comparable.

VERIFIED DATABASE

177

responses

- Recruitment was confirmed by the study team
- 131 fully completed responses became the human reference point
- Used as the trusted benchmark for normal behaviour

LEGACY DATABASE

1,779

responses

- Older, open recruitment with no identity confirmation
- This is where nearly all the suspicious activity sits
- Screened with the same rules, not assumed bad

WHAT WE CONFIRMED FIRST

Response counts matched the expected totals, the same survey questions line up across both databases, and every extract was saved with a checksum so the review can be repeated and audited.

Ten automatic checks on every response

The checks are fixed and rule-based. The same response always gets the same result, with no human judgement in the scoring.

HARD CHECKS · SERIOUS PROBLEMS

- R1** Finished the whole survey faster than any real caregiver
- R2** Raced through the main attitudes section
- R5** Answer pattern is an exact copy of another response
- R8** Family answers contradict each other
- R9** Background details are not possible as given

SOFT CHECKS · WORTH A LOOK

- R3** One section finished unusually fast
- R4** Same answer repeated down a whole block
- R6** Several responses arrived in one rapid burst
- R7** Written comments are near-copies of another response
- R10** Failed the built-in survey fraud question (verified database only)

Each threshold comes from the 131 confirmed-real caregivers, not from a guess. The speed limits are 11.57 minutes for the full survey and 7.85 minutes for the attitudes section.

R10 is a validation signal only. It never decides a response's tier or a payment on its own.

One hard problem, or two soft ones

A single check never condemns a response. Tier 1 needs a hard problem; everything below is graded by how many soft signals stack up.

Tier 1

Confirmed invalid

Any one hard check fired

R1, R2, R5, R8 or R9

Tier 2

High suspicion

Two or more soft checks fired

2+ of R3, R4, R6, R7

Tier 3

Uncertain

Exactly one soft check fired

1 of R3, R4, R6, R7

Tier 4

Pass

Nothing fired at all

No rules triggered


WHY IT IS BUILT THIS WAY

- Tier 1 is the only group that can lose a payment, and even then not automatically.
- Tiers 2, 3 and 4 stay eligible; they differ only in how much checking happens first.

Real caregivers never finish this fast

The speed limit was set by real caregivers, not chosen by us. Not one confirmed-real response fell below it.

- Every confirmed-real caregiver took at least 11.57 minutes.
- In the legacy database, a visible group finishes well under that floor.
- That gap is what the two hard timing checks, R1 and R2, are measuring.
- Because no real caregiver was below the line, these two checks flagged no confirmed-real person at all.




Most hard flags are about speed, not deception

The great majority of serious flags are two timing checks firing together. Contradictory answers are rare, and copied responses never appeared at all.

- 279 responses tripped both timing checks, and for 276 of them that was the only serious problem.
- Only 44 showed contradictory family answers.
- Not one response in either database was an exact copy of another.
- Only 6 combined a contradiction with a speed problem, and those are the strongest cases.

Reading the chart: "dirty_4581" is the legacy database. The diagonal is how often each check fired; off-diagonal cells are two checks firing on the same response.



THE RESULT



Where all 1,956 responses landed

Six responses are confirmed automated. Just over half need a human to look before any money moves.

223

11.4%

Clean pass

No check fired at all.
Pay now.

679

34.7%

Likely real

One soft signal only.
Approve for payment.

1,048

53.6%

Needs review

Hold until a person
adjudicates.

6

0.3%

Confirmed bots

Withhold payment
and block.

THE HEADLINE NUMBER

902 responses are provisionally eligible for a gift card, which is 46% of everything collected. That splits sharply by database: 97% of the verified project clears, against 41% of the legacy project.

Who gets paid, and who does not

Pay everyone the screen cannot fault. Hold, do not refuse, anything uncertain. Refuse only on combined evidence.

SITUATION	ACTION	GIFT CARD
TIER 4 Nothing flagged at all	Pass	Pay
TIER 3 One soft signal only	Conditional pass	Pay
TIER 1 One hard flag on its own	Conditional pass	Pay unless overridden
TIER 2 Two or more soft signals	Review queue	Hold
TIER 1 Two hard flags, speed only	Review queue	Hold
TIER 1 Three or more hard flags	Auto-exclude	Do not pay
TIER 1 Contradiction plus a speed flag	Auto-exclude	Do not pay

Guiding principle: a single rule hit never declares a bot on its own.

Three honest limits on these numbers

The screen is deliberately cautious. It is good at clearing real people and does not claim to catch every bot.

Real people are rarely flagged


No confirmed-real caregiver was flagged by the speed, duplicate or contradiction checks. The soft checks touched at most 4.6% of them, so a small number of genuine caregivers do land in the review queue.

We cannot state a true bot rate

A model suggested most of the legacy database separates from the verified one, but recruitment and survey changes explain part of that. We are not reporting a bot percentage, and some bots almost certainly passed.

The review queue is real work

1,048 responses wait on human adjudication, and 1,043 of them are in the legacy database. Until a person works that list, most payment decisions are provisional rather than final.



Not one confirmed-real caregiver was flagged by the speed, duplicate or contradiction checks. No check uses demographic information as evidence of a bot.